Evidence-Bounded Trace-to-Frontier Replay for Modern NVIDIA GPUs

Chengao Shi
Department of Electronic and Computer Engineering
Hong Kong University of Science and Technology
author email omitted in technical review draft

Technical review draft – October 2026

Abstract

Detailed GPU simulation is valuable precisely where it is expensive: modern kernels combine SIMT scheduling, sectored caches, asynchronous data movement, tensor instructions, and long-latency memory service. We present an evidence-bounded trace-to-frontier replay stack that turns an instrumented NVIDIA execution into a deterministic, architecture-parameterized replay without treating profiler output as a replay input. The stack joins dynamic warp instructions, active and predicate masks, lane-level sectors, placement, and launches with a static SASS-control sidecar; replays four-partition SM issue and dependencies; connects sectored L1, shared L2, and a typed memory service envelope; and emits issue, modeled-completion, observation, and unmodeled-tail frontiers separately. We evaluate the same contract over 286 successful RTX 3090 Ti/SM86 workload captures and 339 successful RTX 5090/SM120 captures, covering 12,392 and 11,558 replayed launches. On the current publication-safe metric subset, global-load sector error is 0.210% (SM86) and 0.010% (SM120); global-L1 lookup-miss error is 9.828% and 6.094%; L2 lookup-miss error is 0.000% and 2.297%; and active IPC error is 14.487% and 9.455%, respectively. We report elapsed IPC below a relaxed 15% bar but do not claim sub-10% closure. Sequential full replay of the 356-root SM86 campaign takes 509.05 s externally (344.819 s of successful-row driver time), compared with 242.530 s for the corresponding non-NCU workload pass; these are campaign-boundary measurements rather than an Accel-Sim speedup. Metric families above the 15% bar, including DRAM-read traffic and some completion frontiers, remain explicit future work. The contribution is the replay contract and full-corpus evidence separation, not recovery of proprietary RTL or a claim that trace-driven/event-driven simulation is new.

1 Introduction

GPU architects need to ask counterfactual questions about cache capacity, latency, scheduling, and memory service before hardware exists. A profiler measures only the installed device, while an execution-driven simulator can change the machine but pays the cost of reconstructing program behavior and microarchitectural state. A trace-driven simulator lies between these extremes: it preserves a recorded dynamic path and spends modeling effort on the proposed architecture. The danger is that a trace can be mistaken for a complete timing history, or that a favorable counter match can hide an unmodeled state variable.

This paper presents a deterministic replay stack designed around that boundary. The input is not an NCU table. It is a dynamic NVBit-style trace and a static SASS control sidecar. The output is not one opaque "cycle" number. It is a set of named frontiers and counters whose causal ownership is explicit. NCU CSV files are joined only after replay, by capture/profile directory and launch identity or ordered fallback, and are retained as independent observations.

Our goal is not to claim a first trace-driven or event-driven GPU simulator; Accel-Sim, MGPUSim/Akita, NVArchSim, and modern GPU-core studies establish those families [5, 7, 8, 4]. Instead, we make four contributions:

1. a concrete dynamic-trace/static-SASS join that preserves warp, lane, sector, placement, and launch identities through issue and memory replay;
2. a deterministic next-event implementation that keeps SM issue, sectored L1/L2 state, memory-service candidates, and wake-up causality in one inspectable path;
3. an evidence contract that separates modeled state from profiler observation and separates issue, completion, observation, and typed-tail frontiers; and

Table 1: Reported contemporary SOTA results from the local literature audit. The values are context, not a cross-paper leaderboard; our launch-level metrics are reported separately in Table 2.

System	Abstraction / target	Reported result	Why it is not pooled here
Modern GPU core	SASS/core, modern NVIDIA	13.45% mean MAPE	Different core model and cohort
Sim-FA	Hopper async pipeline	5.7% MAPE	H800 FA3/TMA scope
GPU-Tile-Sim	Warp-level tiles	1.22–8.71% MAPE	A100/H100 tile workloads
PipeWeave	Analytical + learned	6.1% kernel; 8.5% end-to-end	Prediction, not causal replay
TileSight	First-principles tiles	12.35% single-GPU; 13.52% serving	Tile/network abstraction
LPGSim	Compressed SASS replay	21.4–23.3% cycle error; 9.97× Accel-Sim	Relaxed/local-clock policy

4. a complete-corpus evaluation on Ampere and Blackwell-class capture suites, with an accuracy table restricted to metrics that meet a declared 15% publication threshold and a runtime accounting that distinguishes workload, tracer, profiler, and simulator costs.

The result is a practical architecture-study instrument: fast enough to replay hundreds of real capture directories sequentially, detailed enough to expose sector and readiness mechanisms, and conservative enough to leave unresolved DRAM and completion semantics visible rather than fitting them away.

2 Positioning and research boundary

2.1 Established simulator families

GPGPU-Sim established an execution-driven CUDA/GPU research platform, and Accel-Sim made native NVIDIA instruction traces, correlation, and systematic counter validation central to an open framework [5]. MGPUSim/Akita and NVArchSim demonstrate event/protocol-driven and variable-fidelity organization [7, 8]. Huerta et al. characterize modern NVIDIA scheduler, control-word, register, and dependency behavior [4]. These works rule out "trace-driven," "event-driven," "four-partition," and "validated" as standalone novelty claims.

Recent systems occupy different abstraction points. Sim-FA emphasizes fine-grained asynchronous pipelines; GPU-Tile-Sim/GTSim represents optimized kernels as warp-level tile graphs; PipeWeave combines analytical features and learning; and TileSight provides a first-principles tile/system model. Their goals and denominators differ from instruction-and-sector replay, so we do not pool their errors with ours [10, 3, 9, 6]. Dependency-enriched replay and causal validity are also established concerns; Netrace and CausalSim show why a replay contract must name which dependencies and external responses are preserved [2, 1].

Our claim is therefore a systems integration claim: a trace-to-frontier contract with explicit scope and full-corpus evidence. We do not claim proprietary cache hashes, MSHR counts, scheduler RTL, universal architecture support, or a new general theory of trace sufficiency.

Figure 1 makes the comparison concrete. Accel-Sim owns the public native-SASS validation workflow; MGPUSim/Akita and NVArchSim own event/protocol and industrial variable-fidelity precedents; Huerta et al. set a modern NVIDIA core-fidelity reference; and Sim-FA, GPU-Tile-Sim/GTSim, PipeWeave, and TileSight occupy asynchronous, tile, analytical, and learned abstraction points. Our differentiator is the intersection of native trace identity, sector-aware causal replay, separated frontier semantics, post-replay NCU joins, and complete 286/339-root evidence. The matrix is qualitative and deliberately does not sum heterogeneous capabilities into a single score.

The contemporary reported numbers sharpen the comparison but are not a pooled leaderboard: the MICRO 2025 modern-core model reports 13.45% mean MAPE; Sim-FA reports 5.7% MAPE for its H800 FlashAttention-3 study; GPU-Tile-Sim reports 1.22–8.71% on its A100/H100 tile workloads; PipeWeave reports 6.1% average kernel error and 8.5% end-to-end inference error; and TileSight reports 12.35% pooled single-GPU MAPE and 13.52% serving weighted MAPE. LPGSim reports 21.4–23.3% cycle error across its V100/A100/RTX3090 study and a 9.97× single-thread Accel-Sim speedup. These values use different devices, workloads, abstractions, and denominators; our paper therefore compares capability and scope, then reports our own launch-level MAPE separately.

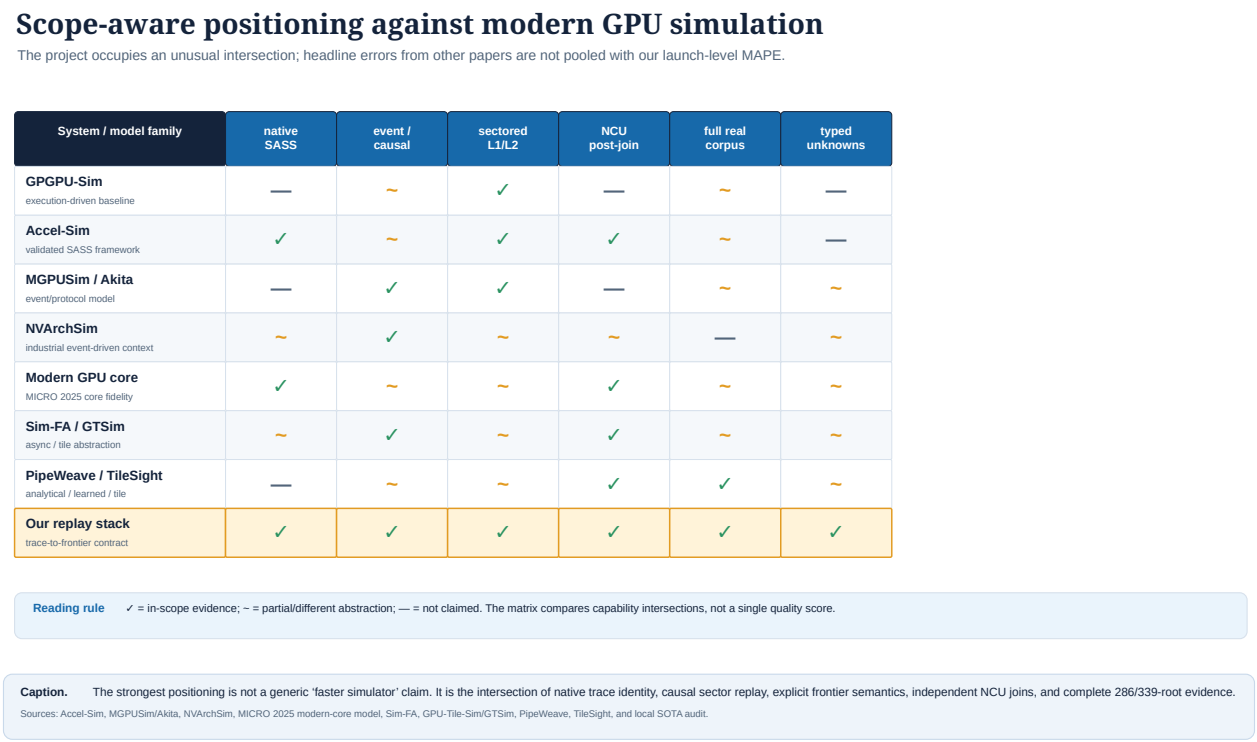


Figure 1: Scope-aware positioning against modern GPU simulation. A checkmark denotes in-scope evidence, a tilde denotes partial or different abstraction, and a dash denotes a capability not claimed. This is a capability intersection, not a scalar ranking of unrelated papers.

3 Replay contract

3.1 Inputs and independent authority

For each capture directory, the replay consumes:

- dynamic records containing PC, warp/block/SM identity, active and predicate masks, per-lane memory addresses, launch geometry, and opcode observations;
- a static SASS sidecar containing stall/yield bits, read/write barriers, wait masks, reuse information, and decoded memory forms; and
- an architecture card containing cache geometry, latencies, service assumptions, and typed unknowns.

The NCU profile is a post-replay authority. The comparator first tries exact capture/profile and launch identity; if compiler-generated kernel names differ, it uses the stable capture directory and launch order. A missing target row is unmatched, not zero error. Scope warnings remain in raw/all-row accounting and are also reported in the secondary like-for-like view.

3.2 State and frontiers

The replay maintains separate state for four SMSP issue partitions, per-warp dependencies and barriers, sectored L1 lines and outstanding fills, shared-L2 lines and source lanes, and an aggregate memory-service envelope. A binary heap orders events by cycle and insertion sequence; idle intervals are skipped without removing causal events.

We expose four frontiers:

1. the last-issue/issue frontier, which identifies the modeled SM work;
2. the modeled completion frontier, which includes declared response and barrier consequences;
3. the observation frontier, which describes how an NCU-shaped endpoint is projected from replay state; and

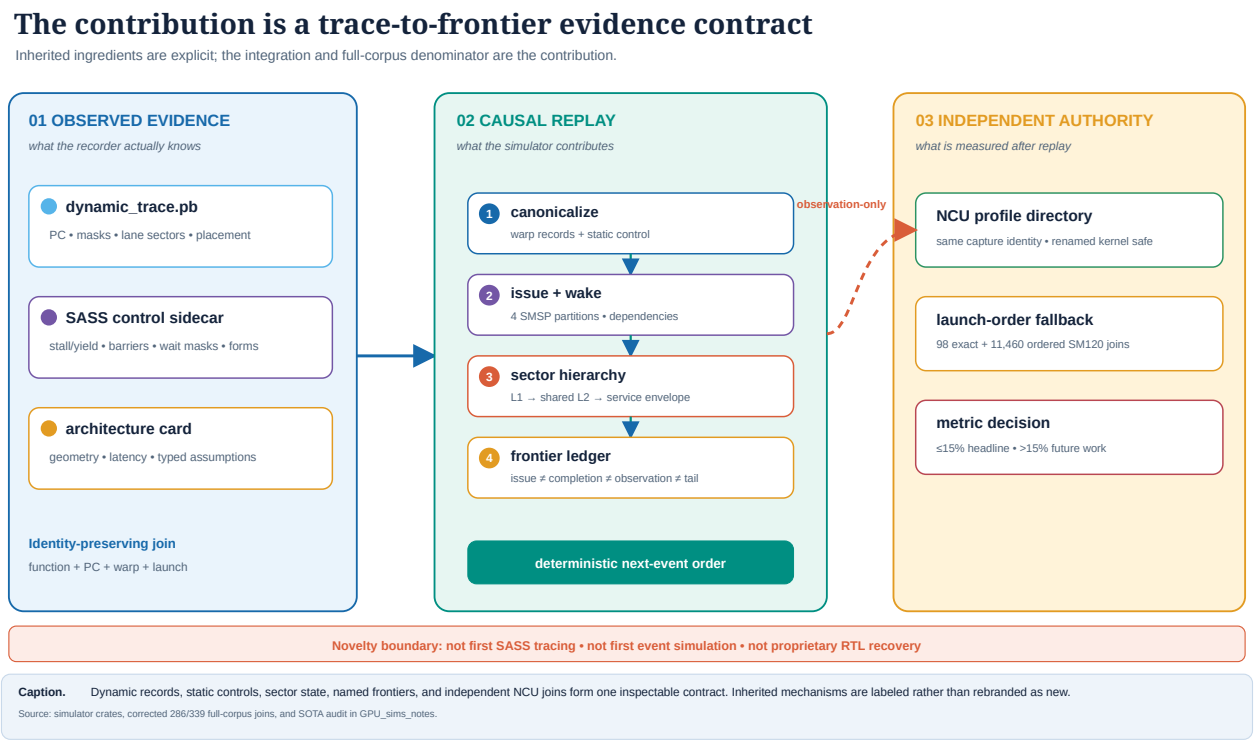


Figure 2: The contribution is a trace-to-frontier evidence contract. The left panel shows what the recorder knows, the middle panel shows the causal replay integration, and the right panel shows the independent NCU authority and launch-scope decision. The bottom boundary states what is inherited rather than claimed as new.

4. typed tails for dirty writeback, coherence, asynchronous visibility, and engine completion that the trace does not establish.

This separation prevents a terminal dirty line or an unobserved asynchronous engine acknowledgement from silently becoming a kernel-cycle claim.

4 SIMT execution and multi-threading strategy

CPU simulation commonly represents each software thread as an independently scheduled instruction stream. Applying that strategy literally to a GPU would create up to 32 host actors per warp, multiply event traffic, and still fail to represent the GPU's lockstep lane semantics. GPU SIMT is hierarchical: threads are grouped into warps, a warp issues one instruction for its active lanes, and multiple resident warps interleave when one warp is stalled on a dependency, barrier, cache miss, or asynchronous operation.

There are three useful modeling choices:

1. **Functional per-thread execution** re-executes lane control flow and values. It can respond to counterfactual timing, but has high host cost and requires a complete functional ISA, memory model, and race semantics.
2. **Warp-level trace replay** stores one dynamic stream per warp, retains active/predicate masks and lane addresses, and schedules warps as the hardware does. It is much cheaper and preserves observed coalescing, divergence, and placement, but cannot regenerate a path that a target intervention would change.
3. **Tile/analytical or learned models** compress many instructions into tile actions or statistical features. They are attractive for throughput and broad design sweeps, but can lose per-lane sector, barrier, and source-qualified cache detail.

Our native SM86/SM120 path uses the second choice. 'sim-sm' is one resumable host actor per SM. Each physical warp slot maps to an SMSP by 'warp_slot mod 4'; each SMSP selects at most one ready warp per issue cycle. The

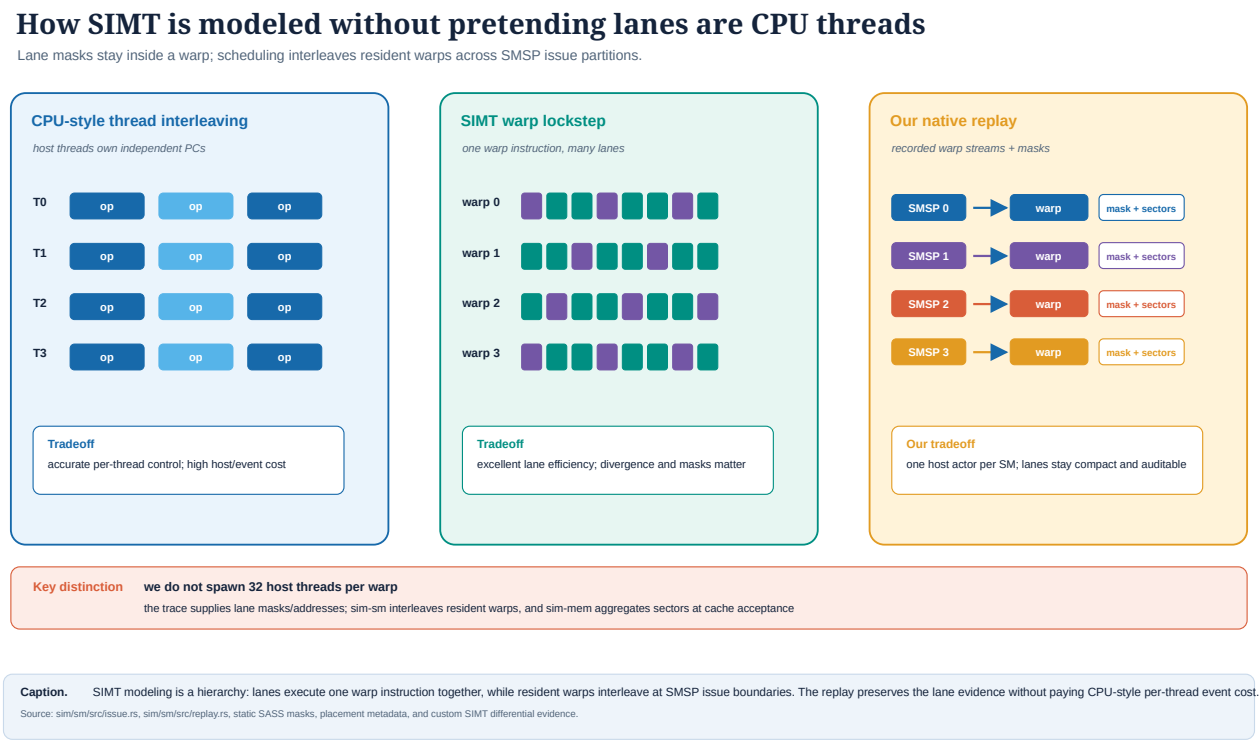


Figure 3: SIMT execution strategy and tradeoffs. CPU-style thread interleaving, warp lockstep, and the replay's one-host-actor-per-SM model are separated. Lane masks stay inside a warp; resident warps interleave at SMSP issue boundaries.

scheduler keeps the held warp while it remains ready, drops the hold on a compiler yield, and otherwise switches to the youngest ready warp. A traced instruction advances exactly one warp stream; its lane mask and sector arena drive memory coalescing, while barriers, dependency waits, residency, and responses determine when that warp can issue again. CTA admission models resident-warp/shared-memory limits and frees resources when a block finishes.

This is not a claim that the host runtime is a proprietary SM scheduler. It is an explicit effective model whose tradeoff is auditable: lane evidence remains available to L1/L2, but the simulator does not pay for 32 independent host threads per warp. The contribution is the identity-preserving connection between SIMT masks, SMSP interleaving, sector state, and named frontiers; the individual ingredients have substantial prior art.

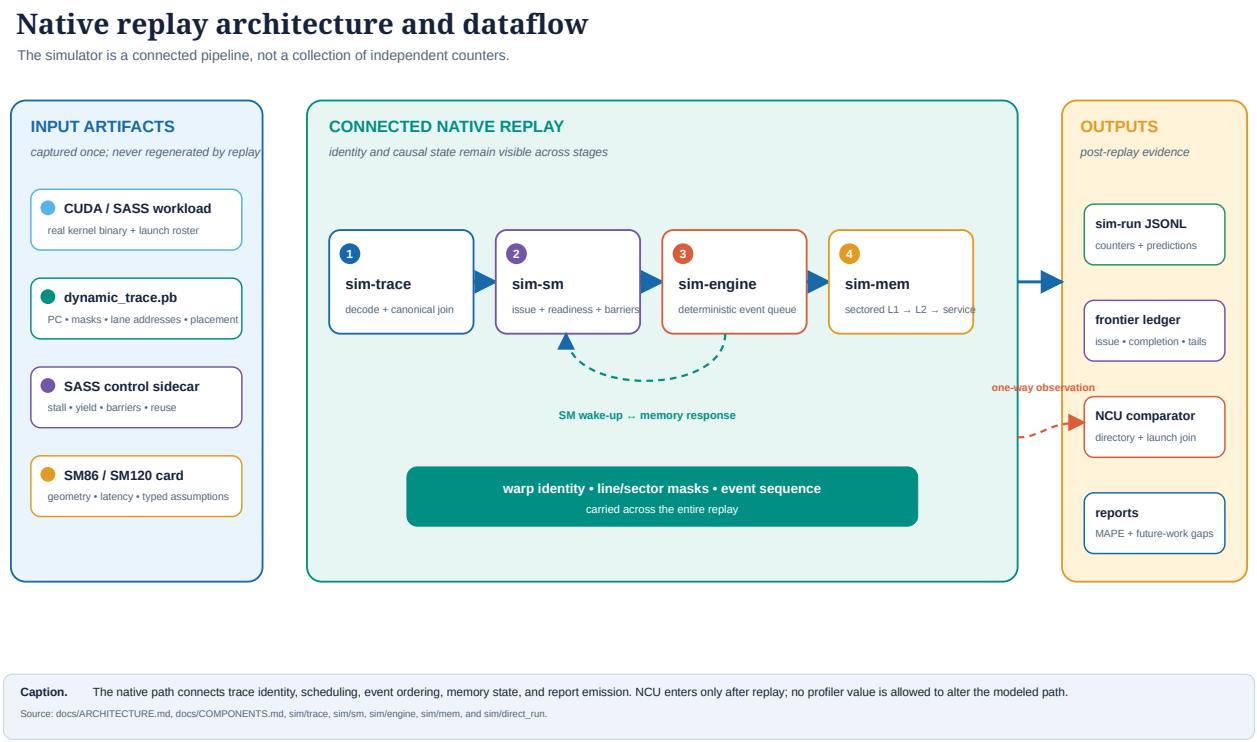
5 Implementation

5.1 Trace reconstruction and issue

The trace reader groups operand observations into one dynamic warp instruction, canonicalizes active/predicate masks, and retains per-lane sectors. The static join is keyed by function identity and PC. The SM model uses four scheduling partitions, compiler-guided stall/yield controls, register/dependency horizons, barrier generations, modeled residency, and response-driven warp wake-up. Every admitted record increments one instruction counter; the replay asserts that the number of issued records equals the records supplied by the capture. This is a replay-integrity invariant, not an independent control-flow prediction.

5.2 Sector-aware memory hierarchy

Lines contain four 32-byte sectors. L1 tracks valid sectors, pending fills, line allocations, and operation-qualified observers. L2 is shared and sectored; reads, stores, atomics, reductions, and asynchronous payload shadows retain distinct source/operation counters. The memory-service layer prices declared demand with a typed aggregate envelope. A prototype DRAM card is enabled for the SM120 full replay, but its bandwidth values are explicitly marked unvalidated; the conservative disabled run is retained.



Caption. The native path connects trace identity, scheduling, event ordering, memory state, and report emission. NCU enters only after replay; no profiler value is allowed to alter the modeled path.

Source: docs/ARCHITECTURE.md, docs/COMPONENTS.md, sim/trace, sim/sm, sim/engine, sim/mem, and sim/direct_run.

Figure 4: Native replay architecture and dataflow. Concrete workload, trace, sidecar, and card artifacts flow through 'sim-trace', 'sim-sm', 'sim-engine', and 'sim-mem' before 'sim-run' emits JSONL counters and frontiers. NCU is a one-way post-replay comparator.

5.3 Metric projections

The simulator emits instruction cardinalities, global load/store sectors, operation-qualified L2 sectors and lookup splits, shared-memory wavefronts, cache-rate projections, IPC projections, and GPC/SM frontier observations. Each prediction carries its source and semantics. For example, a TMA payload sector count is labeled conditional demand rather than an achieved DRAM acknowledgement, and an IPC projection is labeled model-relative rather than a silicon issue-width measurement.

6 Experimental method

6.1 Corpora and replay configuration

The SM86 corpus contains 356 capture roots: 286 successful workload roots, 62 typed mixed-code controls, and 8 replay/input failures. The full replay covers 12,392 launches. The locked reference point is an RTX 3090 Ti at 1350 MHz SM and 10501 MHz memory. The SM120 corpus contains 339 successful roots and 11,558 launches from the RTX 5090 capture suite. The SM120 full replay uses the integrated effective L1/L2 lane; DRAM-disabled and prototype-carded runs are both retained.

All generated replay JSONL artifacts are compressed. Full replay precedes any fast complementary lane. NCU values are never consumed by the simulator or used to select a card parameter.

6.2 Join and scoring rules

For SM86, the latest accepted report has 107 exact identity joins; the publication table uses the metric-specific compatible subset (normally 104– 107 rows). For SM120, compiler-generated names differ more often. We join within the same capture/profile directory and launch order after exact identity; the enabled full artifact yields 98 exact rows and 11,460 ordered-fallback rows, covering all 11,558 replay launches. This is a directory-scoped correspondence, not a string-match assumption.

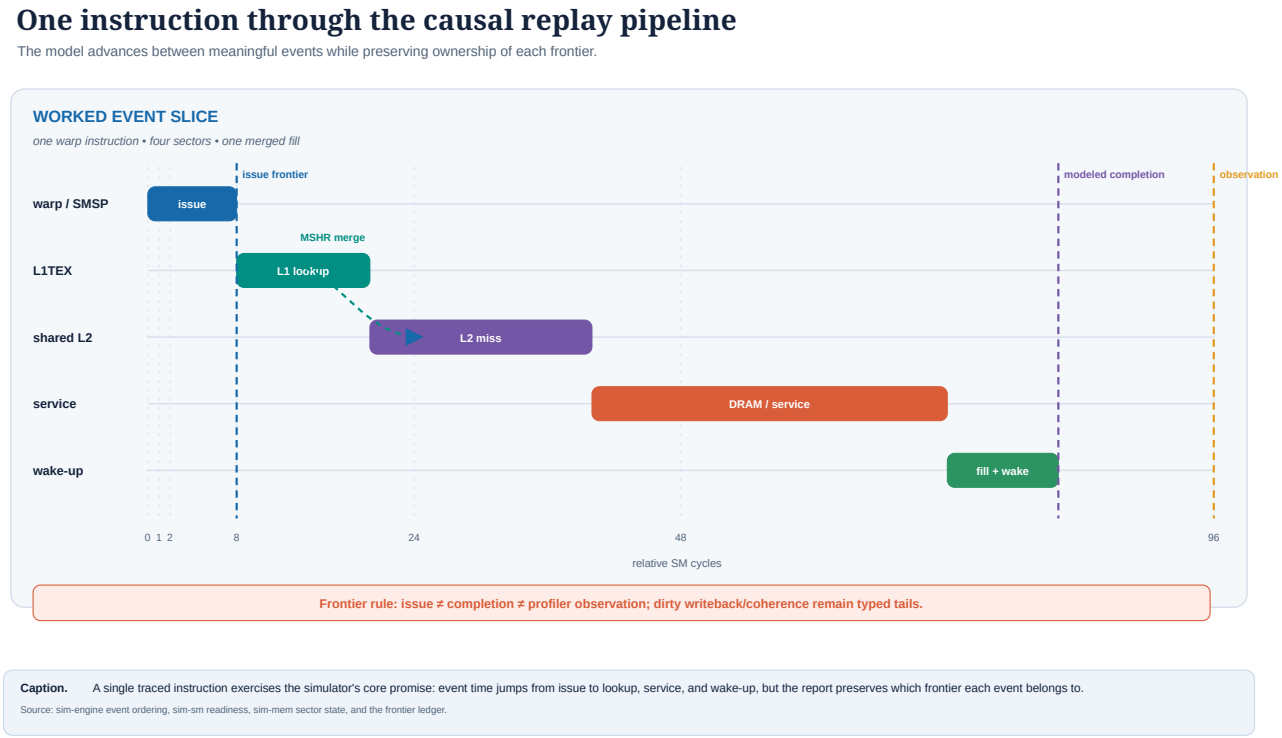


Figure 5: One instruction through the causal replay pipeline. The worked slice separates issue, L1 lookup, shared-L2 miss, service, fill/wake-up, and the three report frontiers; the MSHR merge is explicit rather than hidden in a scalar latency.

The two profile campaigns intentionally have different launch scopes. The SM86 NCU wrapper used '–launch-count 1', so each of its 356 profile directories contains one profiled launch; the 104–107-row metric counts therefore describe all finite target rows present in that one-launch-per-directory campaign, not all 12,392 replay launches. The SM120 wrapper profiles all matching launches, and the corrected 339-directory reduction covers every 11,558 replay launches. This distinction is retained in the denominator rather than hidden by pooling capture counts and launch counts.

For a target y and prediction $\hat{y}$, the row error is

$$e(y, \hat{y}) = \frac{|\hat{y} - y|}{|y|},$$

with zero-target rows handled by exact-zero rules. We report mean APE only for finite nonzero target opportunities and retain raw/all-row scope warnings separately. The paper headline includes metrics at or below 15%; important families above that threshold are future work rather than favorable numbers.

7 Results

7.1 Architectural metric accuracy

Table 2 reports the publication-safe subset. SM86 values are from the H116 full-first report; SM120 values are from the enabled full-only directory/ordered-fallback reduction. Counts are launch-level comparison opportunities, not the number of capture roots.

The table demonstrates two properties. First, sector denominators and operation-qualified lookup-miss surfaces are stable across substantially different targets and workload families. Second, the timing projections are useful but not closed at the preferred 10% threshold on both architectures. We intentionally do not headline DRAM-read traffic, aggregate L2 lookup hits, unified hit rates, or GPC/SM completion families whose current residuals or scope contracts do not meet the declared bar. They remain in the compressed metric reports and in the future-work section.

Table 2: Mean absolute percentage error on complete-corpus replay joins. Only metrics at or below the declared 15% publication bar are headline results. "Active" and "elapsed" IPC use the simulator's model-relative projection semantics.

	SM86 / RTX 3090 Ti			SM120 / RTX 5090		
Metric	MAPE	Rows	Status	MAPE	Rows	Status
Global-load sectors	0.210%	106	report	0.010%	11,558	report
Global-store sectors	0.000%	107	report	0.474%	11,558	report
Global L1 lookup misses	9.828%	106	report	6.094%	11,558	report
L2 TEX lookup misses	0.000%	104	report	2.297%	11,558	report
L2 read lookup misses	0.000%	104	report	0.031%	11,556	report
Shared-load wavefronts	5.623%	107	report	7.250%	11,558	report
Shared-store wavefronts	2.418%	107	report	0.084%	11,558	report
Active IPC	14.487%	107	¡15%	9.455%	11,558	¡10%
Elapsed IPC	12.307%	107	¡15%	13.602%	11,558	¡15%

Table 3: Campaign wall times. NCU and NVBit values are tool-command totals, not bare GPU execution time; no same-host Accel-Sim benchmark is claimed.

Target	Phase	Workloads/roots	Total seconds
RTX 3090 Ti	bare workload, no NCU	286	242.530
RTX 3090 Ti	simulator full external wall	356	509.050
RTX 3090 Ti	simulator driver, successful rows	286	344.819
RTX 3090 Ti	NVBit capture estimate	286	2,640.820
RTX 3090 Ti	NCU profile-command estimate	286	902.797
RTX 5090	simulator full, DRAM disabled	339	511.800
RTX 5090	simulator full, prototype DRAM card	339	512.990
RTX 5090	NVBit capture estimate	339	13,689
RTX 5090	NCU profile-command estimate	339	33,159

7.2 Replay and measurement cost

Table 3 keeps measurement boundaries separate. The SM86 non-NCU workload pass is a sequential process-wall measurement on the locked GPU. NVBit and NCU rows are artifact-timestamp estimates for instrumented and profiled campaign commands. The simulator row is external runner wall time, not just the inner event-loop sum.

The SM86 simulator's successful-row driver sum is 1.42× the bare workload process sum, while the all-root external wall includes typed controls and replay failures. This is strong evidence that the replay is practical for complete-corpus iteration. It is not a controlled claim that the simulator is faster than Accel-Sim: the two systems were not run on the same host, with the same trace, compiler, output, or timeout policy. The 5090 NCU/NVBit totals are also not directly comparable to the 3090 Ti totals because their NCU versions, launch policies, hosts, and profile settings differ.

8 Discussion and limitations

The high-confidence surfaces are those where the trace supplies the relevant denominator and the model preserves sector/operation conservation. L1 lookup misses and L2 source lookup misses therefore meet the publication bar even though some corresponding hit surfaces remain harder: a hit classification is the complement of multiple client and pending-sector states, not a scalar cache capacity law.

The current DRAM-read and completion gaps are not hidden by the table. GCC instruction/constant fetch ownership, maintenance and dirty writeback, exact partition service, asynchronous engine completion, and target-dependent clock crosswalks remain typed or future-work surfaces. The SM120 carded bandwidth values are a prototype hypothesis and must not be read as a locked RTX 5090 measurement.

The trace also fixes the observed dynamic path. If a target intervention changes an atomic order, synchronization arrival, branch, or address calculation, fixed-trace replay answers the timing of the recorded path rather than re-executing the program. This is why the paper reports the trace contract and its boundary alongside accuracy.

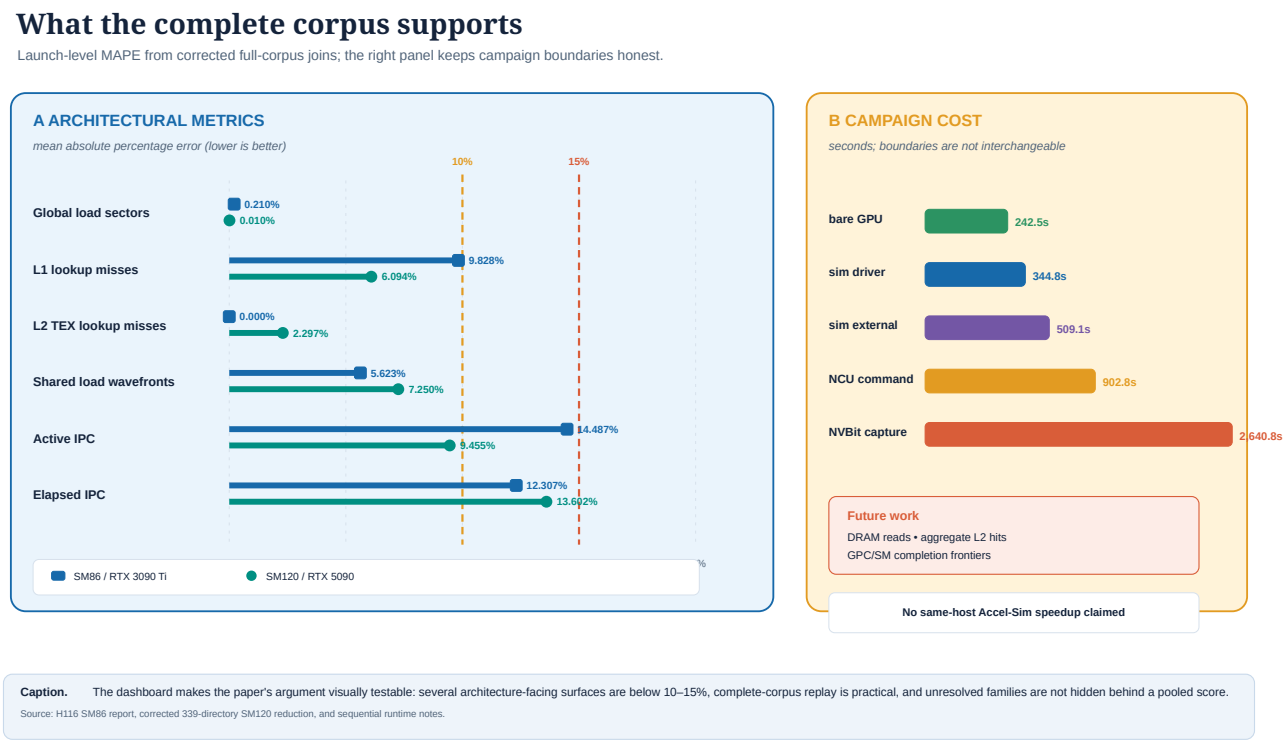


Figure 6: Complete-corpus results and campaign boundaries. Panel (a) shows only architecture-facing metrics at or below the 15% bar, with 10% and 15% thresholds explicit. Panel (b) distinguishes bare GPU, simulator, NCU, and NVBit campaign time; unresolved DRAM and completion families remain future work.

9 Conclusion

We presented an evidence-bounded trace-to-frontier replay stack for modern NVIDIA workload traces. Its practical novelty is the inspectable integration of dynamic lane evidence, static SASS controls, deterministic four-partition issue, sectored L1/L2 state, typed memory-service envelopes, separated frontiers, and an observation-only NCU comparator. Complete replay covers 286 SM86 and 339 SM120 capture roots, and the publication-safe metric subset meets the declared 15% bar with several sectors and lookup-miss surfaces below 10%. Sequential SM86 full replay completes in roughly nine minutes of external runner time, while the corresponding bare workload pass is roughly four minutes; this is a measured campaign comparison, not an unsupported Accel-Sim speedup. DRAM-read and completion families remain future work. The result is a usable, reproducible architecture-study instrument whose claims are scoped to the evidence it actually has.

References

[1] M. Alomar et al. Causalsim: A causal approach to accurate performance modeling. *Proceedings of NSDI*, 2023.

[2] Nathan Binkert et al. Netrace: A trace-based network simulator. In *Proceedings of ISPASS*, 2010.

[3] Y. Ding et al. Gpu-tile-sim: A tile-centric gpu simulation framework for llm hardware-software co-design. *arXiv preprint arXiv:2607.11262*, 2026.

[4] R. Huerta et al. Dissecting and modeling the architecture of modern gpu cores. In *Proceedings of MICRO*, 2025.

[5] Mohamed Khairy, Hany Shen, et al. Accel-Sim: An extensible simulation framework for validated gpu modeling. In *Proceedings of ISCA*, 2020.

[6] Z. Mo et al. Tilesight: A first-principles tile-centric analytical gpu performance model from cores to clusters. *arXiv preprint arXiv:2607.22432*, 2026.

[7] Yifan Sun et al. MGPUSim: Enabling multi-gpu performance modeling and optimization. In *Proceedings of ISCA*, 2019.

[8] Oreste Villa et al. Need for speed: Experiences building a trustworthy system-level gpu simulator. In *Proceedings of HPCA*, 2021.

[9] K. Zhang et al. Pipeweave: Synergizing analytical and learning models for unified gpu performance prediction. *arXiv preprint arXiv:2601.14910*, 2026.

[10] Z. Zhou et al. Sim-fa: A gpgpu simulator framework for fine-grained asynchronous pipeline analysis. *arXiv preprint arXiv:2605.00555*, 2026.